

CentOS7(mini) 安装 Kubernetes 集群 (kubeadm方式)

安装CentOS

1. 安装net-tools

```
1 [root@localhost ~]# yum install -y net-tools
```

1. 关闭firewalld

```
1 [root@localhost ~]# systemctl stop firewalld && systemctl disable firewalld
2 Removed symlink /etc/systemd/system/multi-
  user.target.wants/firewalld.service.
3 Removed symlink /etc/systemd/system/dbus-
  org.fedoraproject.FirewallD1.service.
4 [root@localhost ~]# setenforce 0
5 [root@localhost ~]# sed -i 's/SELINUX=enforcing/SELINUX=disabled/g'
  /etc/selinux/config
```

安装Docker

如今Docker分为了Docker-CE和Docker-EE两个版本，CE为社区版即免费版，EE为企业版即商业版。我们选择使用CE版。

1. 安装yum源工具包

```
1 [root@localhost ~]# yum install -y yum-utils device-mapper-persistent-data
  lvm2
```

1. 下载docker-ce官方的yum源配置文件

```
1 [root@localhost ~]# yum-config-manager --add-repo
  https://download.docker.com/linux/centos/docker-ce.repo
```

1. 禁用docker-ce-edge源配edge是不开发版，不稳定，下载stable版

```
1 yum-config-manager --disable docker-ce-edge
```

1. 更新本地YUM源缓存

```
1 yum makecache fast
```

1. 安装Docker-ce相应版本的

```
1 | yum -y install docker-ce
```

1. 运行hello world

```
1 [root@localhost ~]# systemctl start docker
2 [root@localhost ~]# docker run hello-world
3 Unable to find image 'hello-world:latest' locally
4 latest: Pulling from library/hello-world
5 9a0669468bf7: Pull complete
6 Digest:
   sha256:0e06ef5e1945a718b02a8c319e15bae44f47039005530bc617a5d071190ed3fc
7 Status: Downloaded newer image for hello-world:latest
8
9 Hello from Docker!
10 This message shows that your installation appears to be working correctly.
11
12 To generate this message, Docker took the following steps:
13 1. The Docker client contacted the Docker daemon.
14 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
15 3. The Docker daemon created a new container from that image which runs
   the
16     executable that produces the output you are currently reading.
17 4. The Docker daemon streamed that output to the Docker client, which
   sent it
18     to your terminal.
19
20 To try something more ambitious, you can run an Ubuntu container with:
21 $ docker run -it ubuntu bash
22
23 Share images, automate workflows, and more with a free Docker ID:
24 https://cloud.docker.com/
25
26 For more examples and ideas, visit:
27 https://docs.docker.com/engine/userguide/
```

安装kubelet与kubeadm包

使用kubeadm init命令初始化集群之下载Docker镜像到所有主机的实始化时会下载kubeadm必要的依赖镜像，同时安装etcd,kube-dns,kube-proxy,由于我们GFW防火墙问题我们不能直接访问，因此先通过其它方法下载下面列表中的镜像，然后导入到系统中，再使用kubeadm init来初始化集群

1. 使用[DaoCloud](#)加速器(可以跳过这一步)

```

1 [root@localhost ~]# curl -sSL
  https://get.daocloud.io/daotools/set_mirror.sh | sh -s
  http://0d236e3f.m.daocloud.io
2 docker version >= 1.12
3 {"registry-mirrors": ["http://0d236e3f.m.daocloud.io"]}
4 Success.
5 You need to restart docker to take effect: sudo systemctl restart docker
6 [root@localhost ~]# systemctl restart docker

```

1. 下载镜像,自己通过[Dockerfile](#)到[dockerhub](#)生成对镜像,也可以克隆我的

```

1 images=(kube-controller-manager-amd64 etcd-amd64 k8s-dns-sidecar-amd64
  kube-proxy-amd64 kube-apiserver-amd64 kube-scheduler-amd64 pause-amd64 k8s-
  dns-dnsmasq-nanny-amd64 k8s-dns-kube-dns-amd64)
2 for imageName in ${images[@]} ; do
3     docker pull champly/$imageName
4     docker tag champly/$imageName gcr.io/google_containers/$imageName
5     docker rmi champly/$imageName
6 done

```

1. 修改版本

```

1 docker tag gcr.io/google_containers/etcd-amd64
  gcr.io/google_containers/etcd-amd64:3.0.17 && \
2 docker rmi gcr.io/google_containers/etcd-amd64 && \
3 docker tag gcr.io/google_containers/k8s-dns-dnsmasq-nanny-amd64
  gcr.io/google_containers/k8s-dns-dnsmasq-nanny-amd64:1.14.5 && \
4 docker rmi gcr.io/google_containers/k8s-dns-dnsmasq-nanny-amd64 && \
5 docker tag gcr.io/google_containers/k8s-dns-kube-dns-amd64
  gcr.io/google_containers/k8s-dns-kube-dns-amd64:1.14.5 && \
6 docker rmi gcr.io/google_containers/k8s-dns-kube-dns-amd64 && \
7 docker tag gcr.io/google_containers/k8s-dns-sidecar-amd64
  gcr.io/google_containers/k8s-dns-sidecar-amd64:1.14.2 && \
8 docker rmi gcr.io/google_containers/k8s-dns-sidecar-amd64 && \
9 docker tag gcr.io/google_containers/kube-apiserver-amd64
  gcr.io/google_containers/kube-apiserver-amd64:v1.7.5 && \
10 docker rmi gcr.io/google_containers/kube-apiserver-amd64 && \
11 docker tag gcr.io/google_containers/kube-controller-manager-amd64
  gcr.io/google_containers/kube-controller-manager-amd64:v1.7.5 && \
12 docker rmi gcr.io/google_containers/kube-controller-manager-amd64 && \
13 docker tag gcr.io/google_containers/kube-proxy-amd64
  gcr.io/google_containers/kube-proxy-amd64:v1.6.0 && \
14 docker rmi gcr.io/google_containers/kube-proxy-amd64 && \
15 docker tag gcr.io/google_containers/kube-scheduler-amd64
  gcr.io/google_containers/kube-scheduler-amd64:v1.7.5 && \
16 docker rmi gcr.io/google_containers/kube-scheduler-amd64 && \
17 docker tag gcr.io/google_containers/pause-amd64
  gcr.io/google_containers/pause-amd64:3.0 && \

```

```
18 | docker rmi gcr.io/google_containers/pause-amd64
```

1. 添加阿里源

```
1 [root@localhost ~]# cat >> /etc/yum.repos.d/kubernetes.repo << EOF
2 [kubernetes]
3 name=kubernetes
4 baseurl=https://mirrors.aliyun.com/kubernetes/yum/repos/kubernetes-el7-
  x86_64/
5 enabled=1
6 gpgcheck=0
7 EOF
```

1. 查看kubectl kubelet kubeadm kubernetes-cni列表

```
1 [root@localhost ~]# yum list kubectl kubelet kubeadm kubernetes-cni
2 已加载插件: fastestmirror
3 Loading mirror speeds from cached hostfile
4 * base: mirrors.tuna.tsinghua.edu.cn
5 * extras: mirrors.sohu.com
6 * updates: mirrors.sohu.com
7 可安装的软件包
8 kubeadm.x86_64                                     1.7.5-0
                                                    kubernetes
9 kubectl.x86_64                                     1.7.5-0
                                                    kubernetes
10 kubelet.x86_64                                     1.7.5-0
                                                    kubernetes
11 kubernetes-cni.x86_64                             0.5.1-0
                                                    kubernetes
12 [root@localhost ~]#
```

1. 安装kubectl kubelet kubeadm kubernetes-cni

```
1 [root@localhost ~]# yum install -y kubectl kubelet kubeadm kubernetes-cni
```

修改cgroups

```
1 vi /etc/systemd/system/kubelet.service.d/10-kubeadm.conf
```

```
update KUBELET_CGROUP_ARGS=--cgroup-driver=systemd to KUBELET_CGROUP_ARGS=--
cgroup-driver=cgroupfs
```

修改kubelet中的cAdvisor监控的端口，默认为0改为4194，这样就可以通过浏览器查看kubelet的监控cAdvisor的web页

```
1 [root@kub-master ~]# vi /etc/systemd/system/kubelet.service.d/10-  
  kubeadm.conf
```

```
Environment="KUBELET_CADVISOR_ARGS=--cadvisor-port=4194"
```

启动所有主机上的kubelet服务

```
1 [root@master ~]# systemctl enable kubelet && systemctl start kubelet
```

初始化master master节点上操作

```
1 [root@master ~]# kubeadm reset && kubeadm init --apiserver-advertise-  
  address=192.168.0.100 --kubernetes-version=v1.7.5 --pod-network-  
  cidr=10.200.0.0/16  
2 [preflight] Running pre-flight checks  
3 [reset] Stopping the kubelet service  
4 [reset] Unmounting mounted directories in "/var/lib/kubelet"  
5 [reset] Removing kubernetes-managed containers  
6 [reset] Deleting contents of stateful directories: [/var/lib/kubelet  
  /etc/cni/net.d /var/lib/docker/shim /var/lib/etcd]  
7 [reset] Deleting contents of config directories:  
  [/etc/kubernetes/manifests /etc/kubernetes/pki]  
8 [reset] Deleting files: [/etc/kubernetes/admin.conf  
  /etc/kubernetes/kubelet.conf /etc/kubernetes/controller-manager.conf  
  /etc/kubernetes/scheduler.conf]  
9 [kubeadm] WARNING: kubeadm is in beta, please do not use it for production  
  clusters.  
10 [init] Using Kubernetes version: v1.7.5  
11 [init] Using Authorization modes: [Node RBAC]  
12 [preflight] Running pre-flight checks  
13 [preflight] WARNING: docker version is greater than the most recently  
  validated version. Docker version: 17.09.0-ce. Max validated version: 1.12  
14 [preflight] Starting the kubelet service  
15 [kubeadm] WARNING: starting in 1.8, tokens expire after 24 hours by  
  default (if you require a non-expiring token use --token-ttl 0)  
16 [certificates] Generated CA certificate and key.  
17 [certificates] Generated API server certificate and key.  
18 [certificates] API Server serving cert is signed for DNS names [master  
  kubernetes kubernetes.default kubernetes.default.svc  
  kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 192.168.0.100]  
19 [certificates] Generated API server kubelet client certificate and key.  
20 [certificates] Generated service account token signing key and public key.  
21 [certificates] Generated front-proxy CA certificate and key.  
22 [certificates] Generated front-proxy client certificate and key.  
23 [certificates] Valid certificates and keys now exist in  
  "/etc/kubernetes/pki"  
24 [kubecfg] wrote kubeConfig file to disk: "/etc/kubernetes/admin.conf"  
25 [kubecfg] wrote kubeConfig file to disk: "/etc/kubernetes/kubelet.conf"
```

```
26 [kubeconfig] wrote KubeConfig file to disk: "/etc/kubernetes/controller-
    manager.conf"
27 [kubeconfig] wrote KubeConfig file to disk:
    "/etc/kubernetes/scheduler.conf"
28 [apiclient] Created API client, waiting for the control plane to become
    ready
29 [apiclient] All control plane components are healthy after 34.002949
    seconds
30 [token] Using token: 0696ed.7cd261f787453bd9
31 [apiconfig] Created RBAC rules
32 [addons] Applied essential addon: kube-proxy
33 [addons] Applied essential addon: kube-dns
34
35 Your Kubernetes master has initialized successfully!
36
37 To start using your cluster, you need to run (as a regular user):
38
39     mkdir -p $HOME/.kube
40     sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
41     sudo chown $(id -u):$(id -g) $HOME/.kube/config
42
43 You should now deploy a pod network to the cluster.
44 Run "kubectl apply -f [podnetwork].yaml" with one of the options listed
    at:
45     http://kubernetes.io/docs/admin/addons/
46
47 You can now join any number of machines by running the following on each
    node
48 as root:
49
50     kubeadm join --token 0696ed.7cd261f787453bd9 192.168.0.100:6443
51
52 [root@master ~]#
```

kubeadm join --token 0696ed.7cd261f787453bd9 192.168.0.100:6443 这个一定要记住,以后无法重现,添加节点需要

添加节点

```
1 [root@node1 ~]# kubeadm join --token 0696ed.7cd261f787453bd9
    192.168.0.100:6443
2 [kubeadm] WARNING: kubeadm is in beta, please do not use it for production
    clusters.
3 [preflight] Running pre-flight checks
4 [preflight] WARNING: docker version is greater than the most recently
    validated version. Docker version: 17.09.0-ce. Max validated version: 1.12
5 [preflight] WARNING: kubelet service is not enabled, please run 'systemctl
    enable kubelet.service'
6 [preflight] Starting the kubelet service
```

```

7 [discovery] Trying to connect to API Server "192.168.0.100:6443"
8 [discovery] Created cluster-info discovery client, requesting info from
  "https://192.168.0.100:6443"
9 [discovery] Cluster info signature and contents are valid, will use API
  Server "https://192.168.0.100:6443"
10 [discovery] Successfully established connection with API Server
    "192.168.0.100:6443"
11 [bootstrap] Detected server version: v1.7.10
12 [bootstrap] The server supports the Certificates API
    (certificates.k8s.io/v1beta1)
13 [csr] Created API client to obtain unique certificate for this node,
    generating keys and certificate signing request
14 [csr] Received signed certificate from the API server, generating
    kubeconfig...
15 [kubeconfig] wrote kubeconfig file to disk: "/etc/kubernetes/kubelet.conf"
16
17 Node join complete:
18 * Certificate signing request sent to master and response
19   received.
20 * kubelet informed of new secure connection details.
21
22 Run 'kubectl get nodes' on the master to see this machine join.

```

在master配置kubectl的kubeconfig文件

```

1 [root@master ~]# mkdir -p $HOME/.kube
2 [root@master ~]# cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
3 [root@master ~]# chown $(id -u):$(id -g) $HOME/.kube/config

```

在Master上安装flannel

```

1 docker pull quay.io/coreos/flannel:v0.8.0-amd64
2 kubectl apply -f
  https://raw.githubusercontent.com/coreos/flannel/v0.8.0/Documentation/kube-
  flannel.yml
3 kubectl apply -f
  https://raw.githubusercontent.com/coreos/flannel/v0.8.0/Documentation/kube-
  flannel-rbac.yml

```

查看集群

```

1 [root@master ~]# kubectl get cs
2 NAME                STATUS    MESSAGE              ERROR
3 scheduler            Healthy   ok
4 controller-manager   Healthy   ok
5 etcd-0               Healthy   {"health": "true"}
6 [root@master ~]# kubectl get nodes
7 NAME                STATUS    AGE      VERSION

```

```

8 master    Ready    24m    v1.7.5
9 node1     NotReady 45s    v1.7.5
10 node2    NotReady 7s     v1.7.5
11 [root@master ~]# kubectl get pods --all-namespaces
12 NAMESPACE   NAME                                READY   STATUS
13 kube-system  etcd-master                        1/1     Running
14 kube-system  kube-apiserver-master             1/1     Running
15 kube-system  kube-controller-manager-master    1/1     Running
16 kube-system  kube-dns-2425271678-h48rw         0/3     ImagePullBackOff
17 kube-system  kube-flannel-ds-28n3w             1/2     CrashLoopBackOff
18 kube-system  kube-flannel-ds-ndspr             0/2     ContainerCreating
19 kube-system  kube-flannel-ds-zvx9j             0/2     ContainerCreating
20 kube-system  kube-proxy-qxxzr                  0/1     ImagePullBackOff
21 kube-system  kube-proxy-shkxm                  0/1     ImagePullBackOff
22 kube-system  kube-proxy-vtk52                  0/1     ContainerCreating
23 kube-system  kube-scheduler-master             1/1     Running
24 [root@master ~]#

```

如果出现：The connection to the server localhost:8080 was refused - did you specify the right host or port?

解决办法： 为了使用kubectl访问apiserver，在~/.bash_profile中追加下面的环境变量： export KUBECONFIG=/etc/kubernetes/admin.conf source ~/.bash_profile 重新初始化kubectl